

MLP_Tensorflow

September 16, 2020

0.1 MLP: Initializations

```
[0]: ## MLP: Initializations# https://leonardoaraujosantos.gitbooks.io/  
↪artificial-intelligence/content/multi_layer_perceptron_mnist.html  
# https://github.com/wagonhelm/NaNmnist/blob/master/NaNmnist.ipynb  
# https://github.com/aymericdamien/TensorFlow-Examples/blob/master/notebooks/  
↪3_NeuralNetworks/neural_network.ipynb  
  
from tensorflow.examples.tutorials.mnist import input_data  
import tensorflow as tf  
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
```

WARNING:tensorflow:From <ipython-input-1-0db05428d778>:4: read_data_sets (from tensorflow.contrib.learn.python.learn.datasets.mnist) is deprecated and will be removed in a future version.

Instructions for updating:

Please use alternatives such as official/mnist/dataset.py from tensorflow/models.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/contrib/learn/python/learn/datasets/mnist.py:260: maybe_download (from tensorflow.contrib.learn.python.learn.datasets.base) is deprecated and will be removed in a future version.

Instructions for updating:

Please write your own downloading logic.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/contrib/learn/python/learn/datasets/mnist.py:262: extract_images (from tensorflow.contrib.learn.python.learn.datasets.mnist) is deprecated and will be removed in a future version.

Instructions for updating:

Please use tf.data to implement this functionality.

Extracting MNIST_data/train-images-idx3-ubyte.gz

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/contrib/learn/python/learn/datasets/mnist.py:267: extract_labels (from tensorflow.contrib.learn.python.learn.datasets.mnist) is deprecated and will be removed in a future version.

Instructions for updating:

Please use tf.data to implement this functionality.

Extracting MNIST_data/train-labels-idx1-ubyte.gz

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/contrib/learn/python/learn/datasets/mnist.py:110: dense_to_one_hot (from tensorflow.contrib.learn.python.learn.datasets.mnist) is deprecated and will be removed in a future version.
Instructions for updating:
Please use tf.one_hot on tensors.
Extracting MNIST_data/t10k-images-idx3-ubyte.gz
Extracting MNIST_data/t10k-labels-idx1-ubyte.gz
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/contrib/learn/python/learn/datasets/mnist.py:290: DataSet.__init__ (from tensorflow.contrib.learn.python.learn.datasets.mnist) is deprecated and will be removed in a future version.
Instructions for updating:
Please use alternatives such as official/mnist/dataset.py from tensorflow/models.

```
[0]: # %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, y, y_1, ax, ticks, title, colors=['b']):
    ax.plot(x, y, 'b', label="Train Loss")
    ax.plot(x, y_1, 'r', label="Test Loss")
    if len(x)=Fkee=1:
        plt.legend()
        plt.title(title)
    plt.yticks(ticks)
    fig.canvas.draw()
```

```
[0]: # Network Parameters
n_hidden_1 = 512 # 1st layer number of neurons
n_hidden_2 = 128 # 2nd layer number of neurons
n_input = 784 # MNIST data input (img shape: 28*28)
n_classes = 10 # MNIST total classes (0-9 digits)
```

```
[0]: # x isn't a specific value. It's a placeholder, a value that we'll input when we
    ↪ ask TensorFlow to run a computation.
# We want to be able to input any number of MNIST images, each flattened into a
    ↪ 784-dimensional vector.
# We represent this as a 2-D tensor of floating-point numbers, with a shape X =
    ↪ [None, 784].
# (Here None means that a dimension can be of any length.)
x = tf.placeholder(tf.float32, [None, 784])
y_ = tf.placeholder(tf.float32, [None, 10])
```

```

# keep_prob: we will be using these placeholders when we use dropouts, while
↳testing model
keep_prob = tf.placeholder(tf.float32)
# keep_prob_input: we will be using these placeholders when we use dropouts,
↳while training model
keep_prob_input = tf.placeholder(tf.float32)

```

```

[0]: # Weight initialization

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# https://www.tensorflow.org/api_docs/python/tf/random_normal
# Outputs random values from a normal distribution mean=0 std=1

# If we sample weights from a normal distribution  $N(0, )$  we satisfy this
↳condition with  $=\sqrt{2/(n_i+n_j+1)}$ .
# h1 =>  $=\sqrt{2/(fan\_in+fan\_out+1)} = 0.039 \Rightarrow N(0, ) = N(0, 0.039)$ 
# h2 =>  $=\sqrt{2/(fan\_in+fan\_out+1)} = 0.055 \Rightarrow N(0, ) = N(0, 0.055)$ 
# out =>  $=\sqrt{2/(fan\_in+fan\_out+1)} = 0.120 \Rightarrow N(0, ) = N(0, 0.120)$ 
# SGD: Xavier/Glorot Normal initialization.
weights_sgd = {
    'h1': tf.Variable(tf.random_normal([n_input, n_hidden_1], stddev=0.039,
↳mean=0)), #784x512 #  $\sqrt{2/(784+512)} = 0.039$ 
    'h2': tf.Variable(tf.random_normal([n_hidden_1, n_hidden_2], stddev=0.055,
↳mean=0)), #512x128 #  $\sqrt{2/(512+128)} = 0.055$ 
    'out': tf.Variable(tf.random_normal([n_hidden_2, n_classes], stddev=0.120,
↳mean=0)) #128x10
}

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu lates
# If we sample weights from a normal distribution  $N(0, )$  we satisfy this
↳condition with  $=\sqrt{2/(n_i)}$ .
# h1 =>  $=\sqrt{2/(fan\_in+1)} = 0.062 \Rightarrow N(0, ) = N(0, 0.062)$ 
# h2 =>  $=\sqrt{2/(fan\_in+1)} = 0.125 \Rightarrow N(0, ) = N(0, 0.125)$ 
# out =>  $=\sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0, ) = N(0, 0.120)$ 
# He Normal initialization.
weights_relu = {
    'h1': tf.Variable(tf.random_normal([n_input, n_hidden_1], stddev=0.062,
↳mean=0)), #784x512
    'h2': tf.Variable(tf.random_normal([n_hidden_1, n_hidden_2], stddev=0.125,
↳mean=0)), #512x128
    'out': tf.Variable(tf.random_normal([n_hidden_2, n_classes], stddev=0.120,
↳mean=0)) #128x10
}

```

```

biases = {
    'b1': tf.Variable(tf.random_normal([n_hidden_1])),           #512x1
    'b2': tf.Variable(tf.random_normal([n_hidden_2])),           #128x1
    'out': tf.Variable(tf.random_normal([n_classes]))           #10x1
}

```

```

[0]: # Parameters
training_epochs = 15
learning_rate = 0.001
batch_size = 100
display_step = 1

```

Model 1: input (784) - sigmoid(512) - sigmoid(128) - softmax(output 10)

```

[0]: # https://leonardoaraujosantos.gitbooks.io/artificial-intelligence/content/
    ↪ multi_layer_perceptron_mnist.html
# Create model
def multilayer_perceptron(x, weights, biases):
    # Use tf.matmul instead of "*" because tf.matmul can change it's dimensions
    ↪ on the fly (broadcast)
    print( 'x:', x.get_shape(), 'W[h1]:', weights['h1'].get_shape(), 'b[h1]:',
    ↪ biases['b1'].get_shape())

    # Hidden layer with Sigmoid activation
    layer_1 = tf.add(tf.matmul(x, weights['h1']), biases['b1'])
    ↪ #(x*weights['h1']) + biases['b1']
    layer_1 = tf.nn.sigmoid(layer_1)
    print( 'layer_1:', layer_1.get_shape(), 'W[h2]:', weights['h2'].
    ↪ get_shape(), 'b[h2]:', biases['b2'].get_shape())

    # Hidden layer with Sigmoid activation
    layer_2 = tf.add(tf.matmul(layer_1, weights['h2']), biases['b2'])
    ↪ (layer_1 * weights['h2']) + biases['b2']
    layer_2 = tf.nn.sigmoid(layer_2)
    print( 'layer_2:', layer_2.get_shape(), 'W[out]:', weights['out'].
    ↪ get_shape(), 'b3:', biases['out'].get_shape())

    # Output layer with Sigmoid activation
    out_layer = tf.matmul(layer_2, weights['out']) + biases['out']
    ↪ (layer_2 * weights['out']) + biases['out']
    out_layer = tf.nn.sigmoid(out_layer)
    print('out_layer:', out_layer.get_shape())

    return out_layer

```

___ Model 1 + AdamOptimizer ___

```
[0]: # Since we are using sigmoid activations in hidden layers we will be using
      ↪ weights that are initialized as weights_sgd
y_sgd = multilayer_perceptron(x, weights_sgd, biases)

# https://www.tensorflow.org/api_docs/python/tf/nn/
      ↪ softmax_cross_entropy_with_logits
cost_sgd = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits =
      ↪ y_sgd, labels = y_))

# https://github.com/amitmac/Question-Answering/issues/2
# there are many optimizers available: https://www.tensorflow.org/versions/r1.2/
      ↪ api_guides/python/train#Optimizers
optimizer_adam = tf.train.AdamOptimizer(learning_rate=learning_rate).
      ↪ minimize(cost_sgd)
optimizer_sgdc = tf.train.GradientDescentOptimizer(learning_rate=learning_rate).
      ↪ minimize(cost_sgd)

with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig, ax = plt.subplots(1, 1)
    ax.set_xlabel('epoch') ; ax.set_ylabel('Soft Max Cross Entropy loss')
    xs, ytrs, ytes = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0.
        test_avg_cost = 0.
        total_batch = int(mnist.train.num_examples/batch_size)

        # Loop over all batches
        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)

            # here c: corresponds to the parameter cost_sgd
            # w : correspondse to the parameter weights_sgd
            # c = sess.run() return the cost after every bath during train
            # w = sess.run() return the weights that are modified after every
            ↪ batch through Back prop
            # w is dict w = {'h1': updated h1 weight vector after the current
            ↪ batch,
            #
            # 'h2': updated h2 weight vector after the current
            ↪ batch,
            #
            # 'out': updated output weight vector after the
            ↪ current batch,
            #
            # }
            # you check these w matrix for every iteration, and check whats
            ↪ happening during back prop
            #
```

```

        # note: sess.run() returns parameter values based on the input
    ↪parameters
        # _, c, w = sess.run([optimizer_adam, cost_sgd, weights_sgd]) it
    ↪returns three parameters
        # _, c = sess.run([optimizer_adam, cost_sgd]) it returns two
    ↪parameters
        # _ = sess.run([optimizer_adam]) it returns one parameter (for the
    ↪input optimizer it return none)
        # c = sess.run([cost_sgd]) it returns one parameter (for the input
    ↪cost return error after the batch)

        # feed_dict={x: batch_xs, y_: batch_ys} here x, y_ should be
    ↪placeholders
        # x, y_ are the input parameters on which the models gets trained

        # here we use AdamOptimizer
        _, c, w = sess.run([optimizer_adam, cost_sgd, weights_sgd],
    ↪feed_dict={x: batch_xs, y_: batch_ys})
        train_avg_cost += c / total_batch
        c = sess.run(cost_sgd, feed_dict={x: mnist.test.images, y_: mnist.
    ↪test.labels})
        test_avg_cost += c / total_batch

        xs.append(epoch)
        ytrs.append(train_avg_cost)
        ytes.append(test_avg_cost)
        plt_dynamic(xs, ytrs, ytes, ax, np.arange(1.3, 1.8, step=0.04),
    ↪"input-sigmoid(512)-sigmoid(128)-sigmoid(output)-AdamOptimizer")

        if epoch%display_step == 0:
            print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".
    ↪format(train_avg_cost), "test cost={:.9f}".format(test_avg_cost))
            plt_dynamic(xs, ytrs, ytes, ax, np.arange(1.3, 1.8, step=0.04),
    ↪"input-sigmoid(512)-sigmoid(128)-sigmoid(output)-AdamOptimizer")

        # we are calculating the final accuracy on the test data
        correct_prediction = tf.equal(tf.argmax(y_sgd,1), tf.argmax(y_,1))
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
        print("Accuracy:", accuracy.eval({x: mnist.test.images, y_: mnist.test.
    ↪labels}))

```

```

x: (?, 784) W[h1]: (784, 512) b[h1]: (512,)
layer_1: (?, 512) W[h2]: (512, 128) b[h2]: (128,)
layer_2: (?, 128) W[out]: (128, 10) b3: (10,)
out_layer: (?, 10)
WARNING:tensorflow:From <ipython-input-8-743a1d69c459>:4:
softmax_cross_entropy_with_logits (from tensorflow.python.ops.nn_ops) is

```

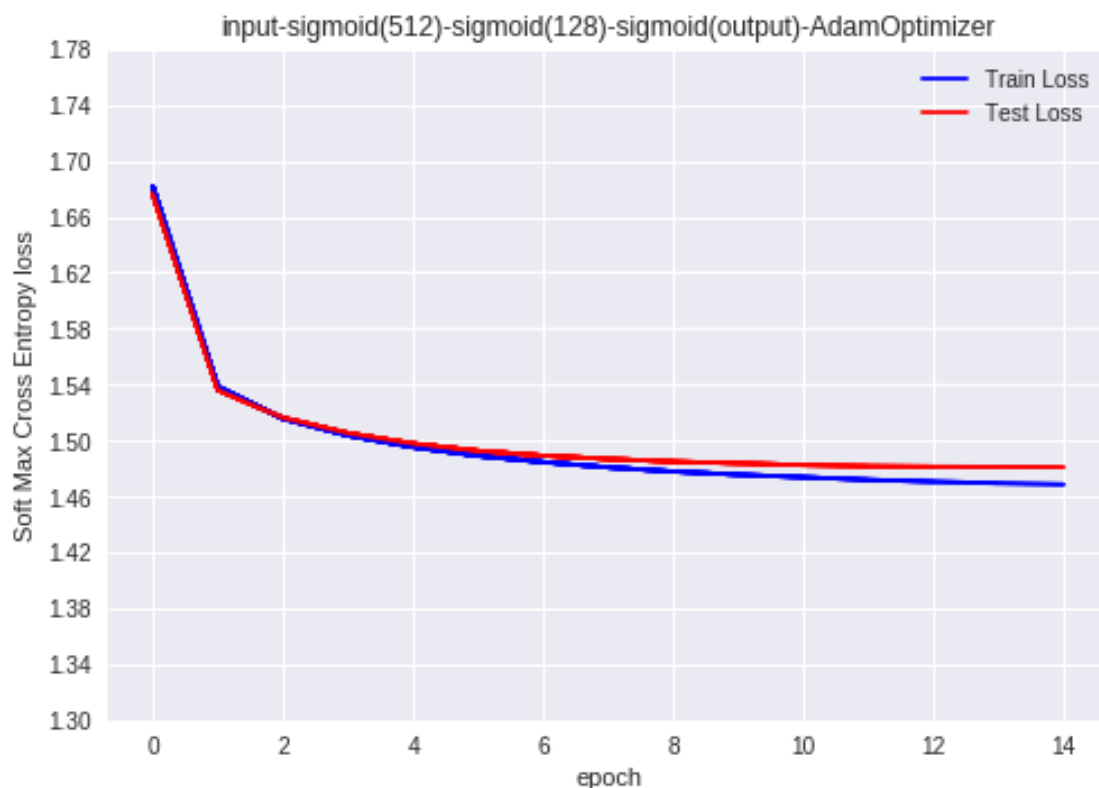
deprecated and will be removed in a future version.

Instructions for updating:

Future major versions of TensorFlow will allow gradients to flow into the labels input on backprop by default.

See `tf.nn.softmax_cross_entropy_with_logits_v2`.

```
Epoch: 0001 train cost=1.681823547 test cost=1.676311211
Epoch: 0002 train cost=1.538490382 test cost=1.536218471
Epoch: 0003 train cost=1.515721546 test cost=1.516237581
Epoch: 0004 train cost=1.503663794 test cost=1.505241357
Epoch: 0005 train cost=1.495330883 test cost=1.497797373
Epoch: 0006 train cost=1.489064606 test cost=1.492622907
Epoch: 0007 train cost=1.484618232 test cost=1.489168651
Epoch: 0008 train cost=1.480839387 test cost=1.486846633
Epoch: 0009 train cost=1.477865693 test cost=1.484905369
Epoch: 0010 train cost=1.475599988 test cost=1.483736856
Epoch: 0011 train cost=1.473791684 test cost=1.482444747
Epoch: 0012 train cost=1.471931230 test cost=1.481658286
Epoch: 0013 train cost=1.470603549 test cost=1.481347544
Epoch: 0014 train cost=1.469495951 test cost=1.481273373
Epoch: 0015 train cost=1.468627278 test cost=1.480902303
Accuracy: 0.9792
```



```
[0]: # Plot weight distribution at the end of training.
```


```
import seaborn as sns
h1_w = w['h1'].flatten().reshape(-1,1)
h2_w = w['h2'].flatten().reshape(-1,1)
out_w = w['out'].flatten().reshape(-1,1)

fig = plt.figure()
plt.subplot(1, 3, 1)
plt.title("Weight matrix")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Weight matrix ")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Weight matrix ")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)
```

___ Model 1 + GradientDescentOptimizer ___

```
[0]: # We can now launch the model in an InteractiveSession

# We first have to create an operation to initialize the variables we created:
# https://github.com/amitmac/Question-Answering/issues/2

# Note: make sure you initialize variables.

with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig, ax = plt.subplots(1, 1)
    ax.set_xlabel('epoch') ; ax.set_ylabel('Soft Max Cross Entropy loss')
    xs, ytrs, ytes = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0.
        test_avg_cost = 0.
        total_batch = int(mnist.train.num_examples/batch_size)

        # Loop over all batches
        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)
```

```

        # here we use GradientDescentOptimizer
        _, c, w = sess.run([optimizer_sgdc, cost_sgd, weights_sgd],
→feed_dict={x: batch_xs, y_: batch_ys})
        train_avg_cost += c / total_batch
        c = sess.run(cost_sgd, feed_dict={x: mnist.test.images, y_: mnist.
→test.labels})
        test_avg_cost += c / total_batch

        xs.append(epoch)
        ytrs.append(train_avg_cost)
        ytes.append(test_avg_cost)
        plt_dynamic(xs, ytrs, ytes, ax, np.arange(2, 2.6, step=0.05),
→"input-sigmoid(512)-sigmoid(128)-sigmoid(output)-GradientDescentOptimizer")

        if epoch%display_step == 0:
            print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".
→format(train_avg_cost), "test cost={:.9f}".format(test_avg_cost))
            plt_dynamic(xs, ytrs, ytes, ax, np.arange(2, 2.6, step=0.05),
→"input-sigmoid(512)-sigmoid(128)-sigmoid(output)-GradientDescentOptimizer")


        # we are calculating the final accuracy on the test data
        correct_prediction = tf.equal(tf.argmax(y_sgd,1), tf.argmax(y_,1))
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
        print("Accuracy:", accuracy.eval({x: mnist.test.images, y_: mnist.test.
→labels}))

```

```

Epoch: 0001 train cost=2.319404572 test cost=2.318047826
Epoch: 0002 train cost=2.316184991 test cost=2.314944556
Epoch: 0003 train cost=2.313415619 test cost=2.312310369
Epoch: 0004 train cost=2.311047141 test cost=2.310042647
Epoch: 0005 train cost=2.309026067 test cost=2.308095952
Epoch: 0006 train cost=2.307287374 test cost=2.306472156
Epoch: 0007 train cost=2.305776415 test cost=2.305014431
Epoch: 0008 train cost=2.304444027 test cost=2.303744715
Epoch: 0009 train cost=2.303249972 test cost=2.302621931
Epoch: 0010 train cost=2.302166392 test cost=2.301561712
Epoch: 0011 train cost=2.301172580 test cost=2.300602047
Epoch: 0012 train cost=2.300254010 test cost=2.299707088
Epoch: 0013 train cost=2.299401874 test cost=2.298892814
Epoch: 0014 train cost=2.298614489 test cost=2.298106769
Epoch: 0015 train cost=2.297889533 test cost=2.297402518
Accuracy: 0.0974

```



```
[0]: h1_w = w['h1'].flatten().reshape(-1,1)
      h2_w = w['h2'].flatten().reshape(-1,1)
      out_w = w['out'].flatten().reshape(-1,1)

      fig = plt.figure()
      plt.subplot(1, 3, 1)
      plt.title("Weight matrix")
      ax = sns.violinplot(y=h1_w,color='b')
      plt.xlabel('Hidden Layer 1')


      plt.subplot(1, 3, 2)
      plt.title("Weight matrix ")
      ax = sns.violinplot(y=h2_w, color='r')
      plt.xlabel('Hidden Layer 2 ')

      plt.subplot(1, 3, 3)
      plt.title("Weight matrix ")
      ax = sns.violinplot(y=out_w,color='y')
      plt.xlabel('Output Layer ')
      plt.show()
```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)

```



Model 2: input (784) - ReLu(512) - ReLu(128) - sigmoid(output 10)

```

[0]: # https://leonardoaraujosantos.gitbooks.io/artificial-intelligence/content/
      ↪ multi_layer_perceptron_mnist.html
# Create model
def multilayer_perceptron_relu(x, weights, biases):
    # Use tf.matmul instead of "*" because tf.matmul can change it's dimensions
    ↪ on the fly (broadcast)
    print( 'x:', x.get_shape(), 'W[h1]:', weights['h1'].get_shape(), 'b[h1]:',
    ↪ biases['b1'].get_shape())

    # Hidden layer with ReLu activation
    # https://en.wikipedia.org/wiki/Rectifier_(neural_networks)
    layer_1 = tf.add(tf.matmul(x, weights['h1']), biases['b1'])
    ↪ #(x*weights['h1']) + biases['b1']

```

```

    layer_1 = tf.nn.relu(layer_1)
    print( 'layer_1:', layer_1.get_shape(), 'W[h2]:', weights['h2'].
    ↳get_shape(), 'b[h2]:', biases['b2'].get_shape())

    # Hidden layer with ReLu activation
    layer_2 = tf.add(tf.matmul(layer_1, weights['h2']), biases['b2']) #
    ↳(layer_1 * weights['h2']) + biases['b2']
    layer_2 = tf.nn.relu(layer_2)
    print( 'layer_2:', layer_2.get_shape(), 'W[out]:', weights['out'].
    ↳get_shape(), 'b3:', biases['out'].get_shape())

    # Output layer with Sigmoid activation
    out_layer = tf.matmul(layer_2, weights['out']) + biases['out'] # (layer_2 *
    ↳weights['out']) + biases['out']
    out_layer = tf.nn.sigmoid(out_layer)
    print('out_layer:',out_layer.get_shape())

    return out_layer

```

___ Input-ReLu(512)-ReLu(128)-sigmoid(output) - AdamOptimizer ___

```

[0]: # Since we are using Relu activations in hidden layers we will be using weights
    ↳that are initialized as weights_relu
yrelu = multilayer_perceptron_relu(x, weights_relu, biases)

# https://www.tensorflow.org/api_docs/python/tf/nn/
    ↳softmax_cross_entropy_with_logits
cost_relu = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits =
    ↳yrelu, labels = y_))
# https://github.com/amitmac/Question-Answering/issues/2
# there are many optimizers available: https://www.tensorflow.org/versions/r1.2/
    ↳api_guides/python/train#Optimizers
optimizer_relu_adam = tf.train.AdamOptimizer(learning_rate=learning_rate).
    ↳minimize(cost_relu)
optimizer_relu_sgdc = tf.train.
    ↳GradientDescentOptimizer(learning_rate=learning_rate).minimize(cost_relu)
with tf.Session() as sess:
    tf.global_variables_initializer().run()

    fig,ax = plt.subplots(1,1)
    ax.set_xlabel('epoch') ; ax.set_ylabel('Soft Max Cross Entropy loss')
    xs, ytrs, ytes = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0.
        test_avg_cost = 0.
        total_batch = int(mnist.train.num_examples/batch_size)

```

```

    # Loop over all batches
    for i in range(total_batch):
        batch_xs, batch_ys = mnist.train.next_batch(batch_size)

        # here we use AdamOptimizer
        _, c, w = sess.run([optimizer_relu_adam, cost_relu, weights_relu],
        ↪feed_dict={x: batch_xs, y_: batch_ys})
        train_avg_cost += c / total_batch
        c = sess.run(cost_relu, feed_dict={x: mnist.test.images, y_: mnist.
        ↪test.labels})
        test_avg_cost += c / total_batch

        xs.append(epoch)
        ytrs.append(train_avg_cost)
        ytes.append(test_avg_cost)
        plt_dynamic(xs, ytrs, ytes, ax, np.arange(1.3, 1.8, step=0.04),
        ↪"input-ReLu(512)-ReLu(128)-sigmoid(output)-AdamOptimizer")

        if epoch%display_step == 0:
            print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".
            ↪format(train_avg_cost), "test cost={:.9f}".format(test_avg_cost))

        # plot final results
        plt_dynamic(xs, ytrs, ytes, ax, np.arange(1.3, 1.8, step=0.04),
        ↪"input-ReLu(512)-ReLu(128)-sigmoid(output)-AdamOptimizer")

        # we are calculating the final accuracy on the test data
        correct_prediction = tf.equal(tf.argmax(yrelu,1), tf.argmax(y_,1))
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
        print("Accuracy:", accuracy.eval({x: mnist.test.images, y_: mnist.test.
        ↪labels}))


```

```

x: (?, 784) W[h1]: (784, 512) b[h1]: (512,)
layer_1: (?, 512) W[h2]: (512, 128) b[h2]: (128,)
layer_2: (?, 128) W[out]: (128, 10) b3: (10,)
out_layer: (?, 10)
Epoch: 0001 train cost=1.562207396 test cost=1.556652784
Epoch: 0002 train cost=1.500463881 test cost=1.501601754
Epoch: 0003 train cost=1.488371294 test cost=1.492026712
Epoch: 0004 train cost=1.482039372 test cost=1.487420058
Epoch: 0005 train cost=1.477783545 test cost=1.484849013
Epoch: 0006 train cost=1.474736516 test cost=1.483547766
Epoch: 0007 train cost=1.472646466 test cost=1.482295828
Epoch: 0008 train cost=1.471696293 test cost=1.482177148
Epoch: 0009 train cost=1.469835149 test cost=1.481312513
Epoch: 0010 train cost=1.469276900 test cost=1.480974347

```

Epoch: 0011 train cost=1.468788507 test cost=1.480446307
Epoch: 0012 train cost=1.468029417 test cost=1.480591303
Epoch: 0013 train cost=1.468475725 test cost=1.480107398
Epoch: 0014 train cost=1.466739849 test cost=1.479299975
Epoch: 0015 train cost=1.466718120 test cost=1.480001667
Accuracy: 0.9785




```
[0]: h1_w = w['h1'].flatten().reshape(-1,1)
      h2_w = w['h2'].flatten().reshape(-1,1)
      out_w = w['out'].flatten().reshape(-1,1)

fig = plt.figure()
plt.subplot(1, 3, 1)
plt.title("Weight matrix")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Weight matrix ")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Weight matrix ")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)
```



Input-ReLu(512)-ReLu(128)-sigmoid(output) - GradientDescentOptimizer

```
[0]: # We can now launch the model in an InteractiveSession

# We first have to create an operation to initialize the variables we created:
# https://github.com/amitmac/Question-Answering/issues/2

# Note: make sure you initialize variables after AdamOptimizer
```



```

with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig, ax = plt.subplots(1,1)
    ax.set_xlabel('epoch') ; ax.set_ylabel('Soft Max Cross Entropy loss')
    xs, ytrs, ytes = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0.
        test_avg_cost = 0.
        total_batch = int(mnist.train.num_examples/batch_size)

        # Loop over all batches
        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)

            # here we use GradientDescentOptimizer
            _, c, w = sess.run([optimizer_relu_sgdc, cost_relu, weights_relu],
→feed_dict={x: batch_xs, y_: batch_ys})
            train_avg_cost += c / total_batch
            c = sess.run(cost_relu, feed_dict={x: mnist.test.images, y_: mnist.
→test.labels})
            test_avg_cost += c / total_batch

        xs.append(epoch)
        ytrs.append(train_avg_cost)
        ytes.append(test_avg_cost)
        plt_dynamic(xs, ytrs, ytes, ax, np.arange(1.5, 2.4, step=0.05),
→"input-ReLu(512)-ReLu(128)-sigmoid(output)-GradientDescentOptimizer")

        if epoch%display_step == 0:
            print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".
→format(train_avg_cost), "test cost={:.9f}".format(test_avg_cost))

        # plot final results
        plt_dynamic(xs, ytrs, ytes, ax, np.arange(1.5, 2.4, step=0.05),
→"input-ReLu(512)-ReLu(128)-sigmoid(output)-GradientDescentOptimizer")

        # we are calculating the final accuracy on the test data
        correct_prediction = tf.equal(tf.argmax(yrelu,1), tf.argmax(y_,1))
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
        print("Accuracy:", accuracy.eval({x: mnist.test.images, y_: mnist.test.
→labels}))

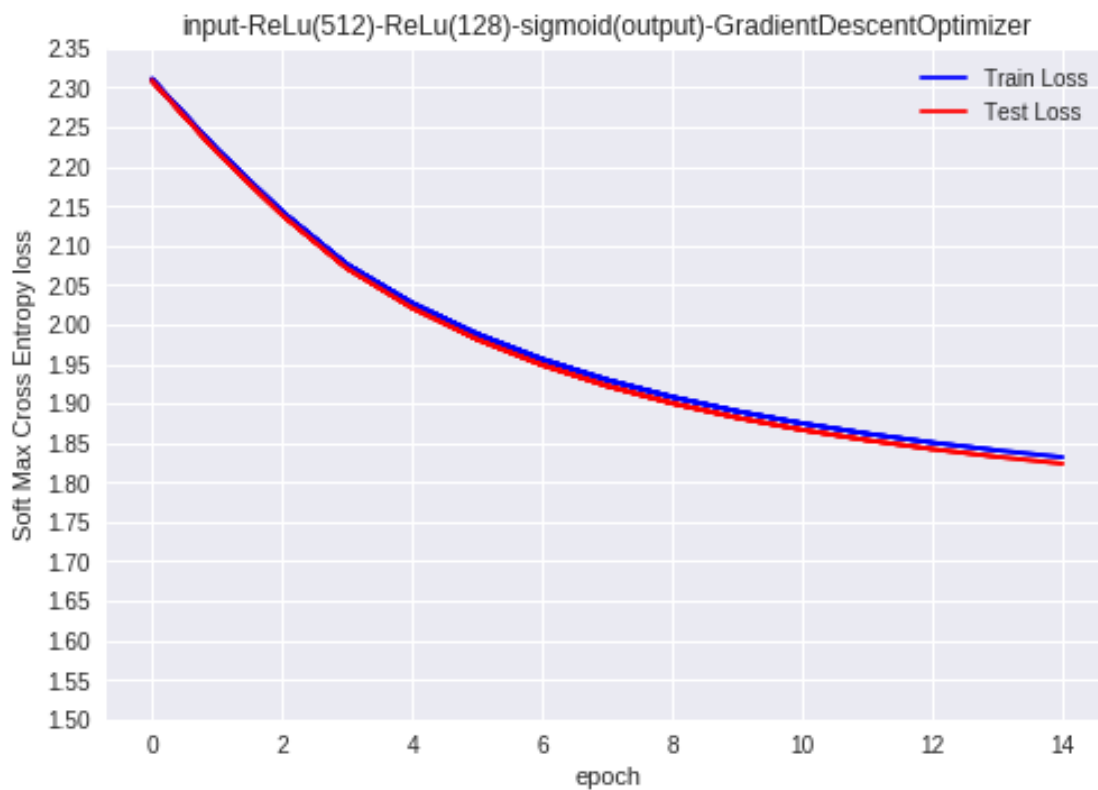
```

```

Epoch: 0001 train cost=2.310429324 test cost=2.309331124
Epoch: 0002 train cost=2.221610156 test cost=2.218818418
Epoch: 0003 train cost=2.142771041 test cost=2.139230260
Epoch: 0004 train cost=2.075905863 test cost=2.071342851

```

Epoch: 0005 train cost=2.027072112 test cost=2.021270330
Epoch: 0006 train cost=1.987932869 test cost=1.981295761
Epoch: 0007 train cost=1.956008127 test cost=1.948708092
Epoch: 0008 train cost=1.929796880 test cost=1.922126883
Epoch: 0009 train cost=1.908102709 test cost=1.900147488
Epoch: 0010 train cost=1.889981364 test cost=1.881854080
Epoch: 0011 train cost=1.874673119 test cost=1.866476579
Epoch: 0012 train cost=1.861609627 test cost=1.853379037
Epoch: 0013 train cost=1.850336982 test cost=1.842097878
Epoch: 0014 train cost=1.840515356 test cost=1.832302262
Epoch: 0015 train cost=1.831871038 test cost=1.823694615
Accuracy: 0.1066



```
[0]: h1_w = w['h1'].flatten().reshape(-1,1)
      h2_w = w['h2'].flatten().reshape(-1,1)
      out_w = w['out'].flatten().reshape(-1,1)


      fig = plt.figure()
      plt.subplot(1, 3, 1)
      plt.title("Weight matrix")
      ax = sns.violinplot(y=h1_w,color='b')
```

```
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Weight matrix ")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Weight matrix ")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
violin_data = remove_na(group_data)



Model 3: Input - Sigmoid(BatchNormalization(512)) - Sigmoid(BatchNormalization(128))- Sigmoid(output)

```

[0]: # https://www.tensorflow.org/api\_docs/python/tf/nn/batch\_normalization
# https://r2rt.com/implementing-batch-normalization-in-tensorflow.html
epsilon = 1e-3
def multilayer_perceptron_batch(x, weights, biases):
    # Use tf.matmul instead of "*" because tf.matmul can change it's dimensions
    ↪ on the fly (broadcast)
    print( 'x:', x.get_shape(), 'W[h1]:', weights['h1'].get_shape(), 'b[h1]:',
    ↪ biases['b1'].get_shape())

    #####
    # Hidden layer with Sigmoid activation and batch normalization
    layer_1 = tf.add(tf.matmul(x, weights['h1']), biases['b1'])
    ↪ #(x*weights['h1']) + biases['b1']

    # https://www.tensorflow.org/api\_docs/python/tf/nn/moments
    # Calculate the mean and variance of x.
    batch_mean_1, batch_var_1 = tf.nn.moments(layer_1, [0])

    scale_1 = tf.Variable(tf.ones([n_hidden_1]))
    beta_1 = tf.Variable(tf.zeros([n_hidden_1]))

    # https://www.tensorflow.org/api\_docs/python/tf/nn/batch\_normalization
    layer_1 = tf.nn.batch_normalization(layer_1, batch_mean_1, batch_var_1,
    ↪ beta_1, scale_1, epsilon)
    layer_1 = tf.nn.sigmoid(layer_1)

    print( 'layer_1:', layer_1.get_shape(), 'W[h2]:', weights['h2'].
    ↪ get_shape(), 'b[h2]:', biases['b2'].get_shape())

    ↪
    ↪ #####

    # Hidden layer with Sigmoid activation and batch normalization
    layer_2 = tf.add(tf.matmul(layer_1, weights['h2']), biases['b2'])
    ↪ (layer_1 * weights['h2']) + biases['b2']

    # https://www.tensorflow.org/api\_docs/python/tf/nn/moments
    # Calculate the mean and variance of x.
    batch_mean_2, batch_var_2 = tf.nn.moments(layer_2, [0])

    scale_2 = tf.Variable(tf.ones([n_hidden_2]))
    beta_2 = tf.Variable(tf.zeros([n_hidden_2]))

    layer_2 = tf.nn.batch_normalization(layer_2, batch_mean_2, batch_var_2,
    ↪ beta_2, scale_2, epsilon)
    layer_2 = tf.nn.sigmoid(layer_2)

```

```

    print( 'layer_2:', layer_2.get_shape(), 'W[out]:', weights['out'].
    ↳get_shape(), 'b3:', biases['out'].get_shape())

    ↳
    ↳#####

    # output layer with Sigmoid activation
    out_layer = tf.matmul(layer_2, weights['out']) + biases['out'] # (layer_2 *
    ↳weights['out']) + biases['out']
    out_layer = tf.nn.sigmoid(out_layer)
    print('out_layer:',out_layer.get_shape())

    return out_layer

```

___ Model 3+ AdamOptimizer ___

```

[0]: # Since we are using sigmoid activations in hidden layers we will be using
    ↳weights that are initialized as weights_sgd
ybatch = multilayer_perceptron_batch(x, weights_sgd, biases)

# https://www.tensorflow.org/api_docs/python/tf/nn/
    ↳softmax_cross_entropy_with_logits
cost_batch = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits =
    ↳ybatch, labels = y_))
# https://github.com/amitmac/Question-Answering/issues/2
# there are many optimizers available: https://www.tensorflow.org/versions/r1.2/
    ↳api_guides/python/train#Optimizers
optimizer_batch_adam = tf.train.AdamOptimizer(learning_rate=learning_rate).
    ↳minimize(cost_batch)
optimizer_batch_sgdc = tf.train.
    ↳GradientDescentOptimizer(learning_rate=learning_rate).minimize(cost_batch)
with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig,ax = plt.subplots(1,1)
    ax.set_xlabel('epoch') ; ax.set_ylabel('Soft Max Cross Entropy loss')
    xs, ytrs, ytes = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0.
        test_avg_cost = 0.
        total_batch = int(mnist.train.num_examples/batch_size)

        # Loop over all batches
        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)

            # here we use AdamOptimizer

```

```

        _, c, w = sess.run([optimizer_batch_adam, cost_batch, weights_sgd],  

↪feed_dict={x: batch_xs, y_: batch_ys})  

        train_avg_cost += c / total_batch  

        c = sess.run(cost_batch, feed_dict={x: mnist.test.images, y_: mnist.  

↪test.labels})  

        test_avg_cost += c / total_batch  
  

        xs.append(epoch)  

        ytrs.append(train_avg_cost)  

        ytes.append(test_avg_cost)  

        plt_dynamic(xs, ytrs, ytes, ax, np.arange(1.3, 1.8, step=0.04),  

↪"input-Sigmoid(BN(512))-Sigmoid(BN(128))-Sigmoid(output)-AdamOptimizer")  
  

        if epoch%display_step == 0:  

            print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".  

↪format(train_avg_cost), "test cost={:.9f}".format(test_avg_cost))  
  

        # plot final results  

        plt_dynamic(xs, ytrs, ytes, ax, np.arange(1.3, 1.8, step=0.04),  

↪"input-Sigmoid(BN(512))-Sigmoid(BN(128))-Sigmoid(output)-AdamOptimizer")  
  

        # we are calculating the final accuracy on the test data  

        correct_prediction = tf.equal(tf.argmax(ybatch,1), tf.argmax(y_,1))  

        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))  

        print("Accuracy:", accuracy.eval({x: mnist.test.images, y_: mnist.test.  

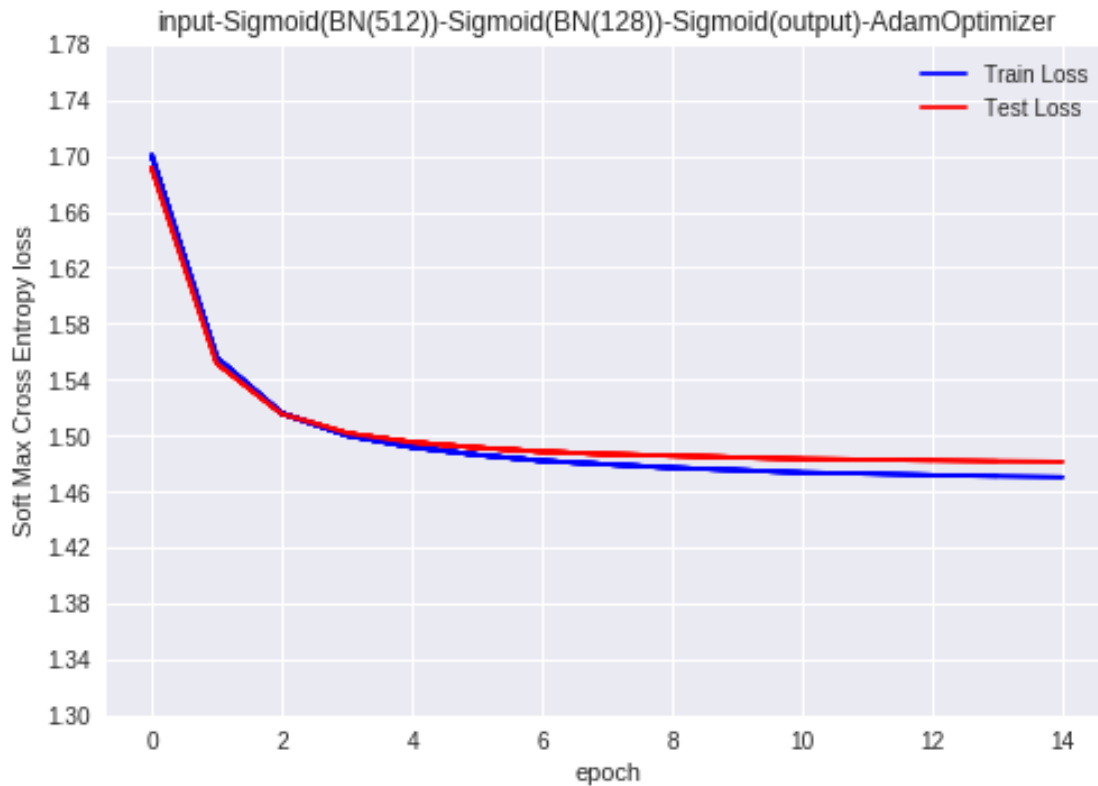
↪labels}))

```

```

x: (?, 784) W[h1]: (784, 512) b[h1]: (512,)
layer_1: (?, 512) W[h2]: (512, 128) b[h2]: (128,)
layer_2: (?, 128) W[out]: (128, 10) b3: (10,)
out_layer: (?, 10)
Epoch: 0001 train cost=1.700068553 test cost=1.691480070
Epoch: 0002 train cost=1.555593932 test cost=1.551793572
Epoch: 0003 train cost=1.515948518 test cost=1.515595129
Epoch: 0004 train cost=1.500090736 test cost=1.501608725
Epoch: 0005 train cost=1.491691740 test cost=1.495176234
Epoch: 0006 train cost=1.486215817 test cost=1.491377480
Epoch: 0007 train cost=1.481944979 test cost=1.488477005
Epoch: 0008 train cost=1.479413016 test cost=1.486683784
Epoch: 0009 train cost=1.476976019 test cost=1.485592169
Epoch: 0010 train cost=1.475328626 test cost=1.484361933
Epoch: 0011 train cost=1.473508603 test cost=1.483245888
Epoch: 0012 train cost=1.472790788 test cost=1.482792196
Epoch: 0013 train cost=1.471692099 test cost=1.482280265
Epoch: 0014 train cost=1.470980778 test cost=1.481615747
Epoch: 0015 train cost=1.470204555 test cost=1.481264244
Accuracy: 0.9811

```



```
[0]: h1_w = w['h1'].flatten().reshape(-1,1)
      h2_w = w['h2'].flatten().reshape(-1,1)
      out_w = w['out'].flatten().reshape(-1,1)

      fig = plt.figure()
      plt.subplot(1, 3, 1)
      plt.title("Weight matrix")
      ax = sns.violinplot(y=h1_w,color='b')
      plt.xlabel('Hidden Layer 1')

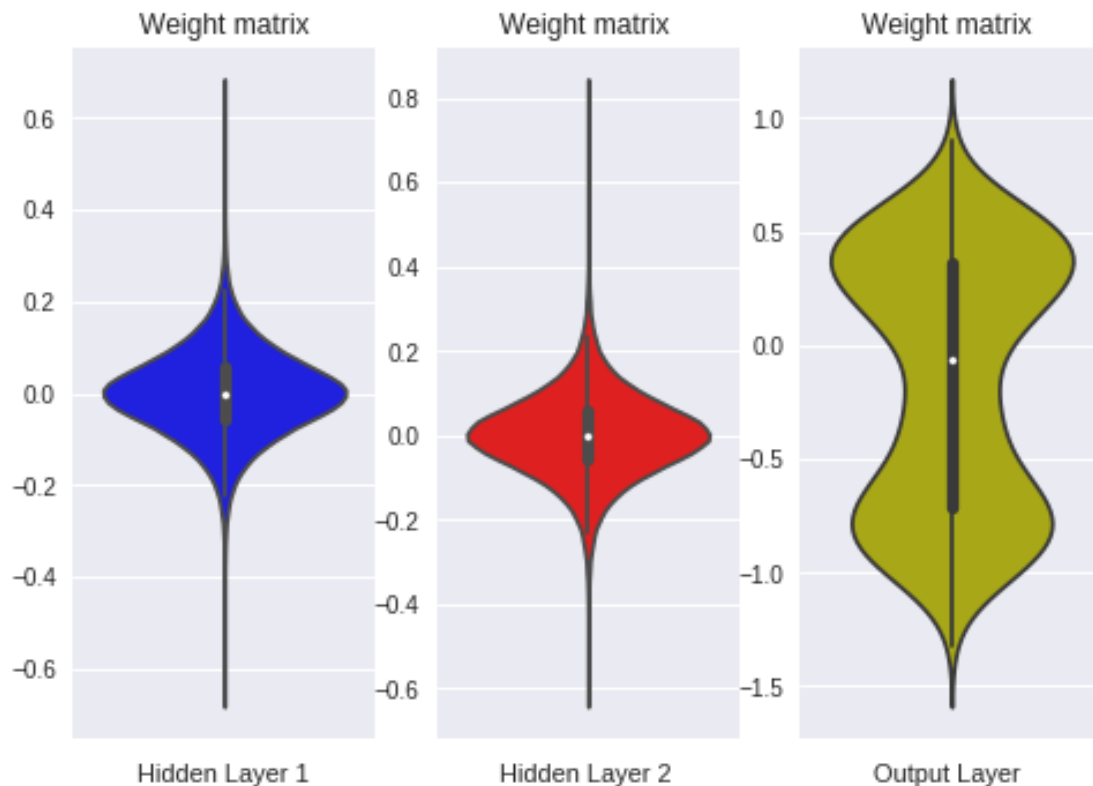
      plt.subplot(1, 3, 2)
      plt.title("Weight matrix ")
      ax = sns.violinplot(y=h2_w, color='r')
      plt.xlabel('Hidden Layer 2 ')

      plt.subplot(1, 3, 3)
      plt.title("Weight matrix ")
      ax = sns.violinplot(y=out_w,color='y')
      plt.xlabel('Output Layer ')
      plt.show()
```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)

```



___ Model 3 + GradientDescentOptimizer ___

```

[0]: with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig, ax = plt.subplots(1, 1)
    ax.set_xlabel('epoch') ; ax.set_ylabel('Soft Max Cross Entropy loss')
    xs, ytrs, ytes = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0.
        test_avg_cost = 0.
        total_batch = int(mnist.train.num_examples/batch_size)

        # Loop over all batches
        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)

```



```

        # here we use GradientDescentOptimizer
        _, c, w = sess.run([optimizer_batch_sgdc, cost_batch, weights_sgd],
→feed_dict={x: batch_xs, y_: batch_ys})
        train_avg_cost += c / total_batch
        c = sess.run(cost_batch, feed_dict={x: mnist.test.images, y_: mnist.
→test.labels})
        test_avg_cost += c / total_batch

        xs.append(epoch)
        ytrs.append(train_avg_cost)
        ytes.append(test_avg_cost)
        plt_dynamic(xs, ytrs, ytes, ax, np.arange(1.5, 2.4, step=0.05),
→"input-Sigmoid(BN(512))-Sigmoid(BN(128))-Sigmoid(output)-GradientDescentOptimizer")

        if epoch%display_step == 0:
            print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".
→format(train_avg_cost), "test cost={:.9f}".format(test_avg_cost))

        # plot final results
        plt_dynamic(xs, ytrs, ytes, ax, np.arange(1.5, 2.4, step=0.05),
→"input-Sigmoid(BN(512))-Sigmoid(BN(128))-Sigmoid(output)-GradientDescentOptimizer")


        # we are calculating the final accuracy on the test data
        correct_prediction = tf.equal(tf.argmax(ybatch,1), tf.argmax(y_,1))
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
        print("Accuracy:", accuracy.eval({x: mnist.test.images, y_: mnist.test.
→labels}))

```

```

Epoch: 0001 train cost=2.312056991 test cost=2.312045197
Epoch: 0002 train cost=2.286723345 test cost=2.286289722
Epoch: 0003 train cost=2.262095340 test cost=2.261143614
Epoch: 0004 train cost=2.237721429 test cost=2.236359837
Epoch: 0005 train cost=2.214522683 test cost=2.212769454
Epoch: 0006 train cost=2.193412623 test cost=2.191231863
Epoch: 0007 train cost=2.174389078 test cost=2.171921795
Epoch: 0008 train cost=2.157141389 test cost=2.154426692
Epoch: 0009 train cost=2.141131058 test cost=2.138286440
Epoch: 0010 train cost=2.126325870 test cost=2.123236751
Epoch: 0011 train cost=2.112224244 test cost=2.109058378
Epoch: 0012 train cost=2.099121217 test cost=2.095732109
Epoch: 0013 train cost=2.086694297 test cost=2.083267775
Epoch: 0014 train cost=2.075242446 test cost=2.071594143
Epoch: 0015 train cost=2.064474456 test cost=2.060690631
Accuracy: 0.125

```



```
[0]: h1_w = w['h1'].flatten().reshape(-1,1)
      h2_w = w['h2'].flatten().reshape(-1,1)
      out_w = w['out'].flatten().reshape(-1,1)

      fig = plt.figure()
      plt.subplot(1, 3, 1)
      plt.title("Weight matrix")
      ax = sns.violinplot(y=h1_w,color='b')
      plt.xlabel('Hidden Layer 1')


      plt.subplot(1, 3, 2)
      plt.title("Weight matrix ")
      ax = sns.violinplot(y=h2_w, color='r')
      plt.xlabel('Hidden Layer 2 ')

      plt.subplot(1, 3, 3)
      plt.title("Weight matrix ")
      ax = sns.violinplot(y=out_w,color='y')
      plt.xlabel('Output Layer ')
      plt.show()
```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)

```



Model 4: Input - ReLu(512) - Dropout - ReLu(128)- Dropout -Sigmoid(output)

```

[0]: # https://leonardoaraujosantos.gitbooks.io/artificial-intelligence/content/
      ↳ multi_layer_perceptron_mnist.html
      # Create model
      def multilayer_perceptron_dropout(x, weights, biases):
          # Use tf.matmul instead of "*" because tf.matmul can change it's dimensions
          ↳ on the fly (broadcast)
          print( 'x:', x.get_shape(), 'W[h1]:', weights['h1'].get_shape(), 'b[h1]:',
          ↳ biases['b1'].get_shape())
          # we are adding a drop out layer after input layers with parameter
          ↳ keep_prob_input

          # Hidden layer with ReLu activation

```

```

    layer_1 = tf.add(tf.matmul(x, weights['h1']), biases['b1'])
    ↪ # (x * weights['h1']) + biases['b1']
    layer_1 = tf.nn.relu(layer_1)
    # we are adding a drop out layer after the first hidden layer with parameter
    ↪ keep_prob
    layer_1_drop = tf.nn.dropout(layer_1, keep_prob)

    print('layer_1:', layer_1.get_shape(), 'W[h2]:', weights['h2'].
    ↪ get_shape(), 'b[h2]:', biases['b2'].get_shape())

    # Hidden layer with ReLu activation
    layer_2 = tf.add(tf.matmul(layer_1_drop, weights['h2']), biases['b2']) #
    ↪ (layer_1 * weights['h2']) + biases['b2']
    layer_2 = tf.nn.relu(layer_2)
    # we are adding a drop out layer after the first hidden layer with parameter
    ↪ keep_prob
    layer_2_drop = tf.nn.dropout(layer_2, keep_prob)
    print('layer_2:', layer_2.get_shape(), 'W[out]:', weights['out'].
    ↪ get_shape(), 'b3:', biases['out'].get_shape())

    # Output layer with Sigmoid activation
    out_layer = tf.matmul(layer_2_drop, weights['out']) + biases['out'] #
    ↪ (layer_2 * weights['out']) + biases['out']
    out_layer = tf.nn.sigmoid(out_layer)
    print('out_layer:', out_layer.get_shape())

    return out_layer

```

___ Model 4 + AdamOptimizer ___

```

[0]: # Since we are using Relu activations in hidden layers we will be using weights
    ↪ that are initialized as weights_relu
    ydrop = multilayer_perceptron_dropout(x, weights_relu, biases)

    # https://www.tensorflow.org/api_docs/python/tf/nn/
    ↪ softmax_cross_entropy_with_logits
    cost_drop = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits =
    ↪ ydrop, labels = y_))
    # https://github.com/amitmac/Question-Answering/issues/2
    # there are many optimizers available: https://www.tensorflow.org/versions/r1.2/
    ↪ api_guides/python/train#Optimizers
    optimizer_drop_adam = tf.train.AdamOptimizer(learning_rate=learning_rate).
    ↪ minimize(cost_drop)
    optimizer_drop_sgdc = tf.train.
    ↪ GradientDescentOptimizer(learning_rate=learning_rate).minimize(cost_drop)

```

```

with tf.Session() as sess:
    tf.global_variables_initializer().run()
    fig,ax = plt.subplots(1,1)
    ax.set_xlabel('epoch') ; ax.set_ylabel('Soft Max Cross Entropy loss')
    xs, ytrs, ytes = [], [], []
    for epoch in range(training_epochs):
        train_avg_cost = 0.
        test_avg_cost = 0.
        total_batch = int(mnist.train.num_examples/batch_size)

        # Loop over all batches
        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)

            # here we use AdamOptimizer
            _, c, w = sess.run([optimizer_drop_adam, cost_drop, weights_relu],
↪feed_dict={x: batch_xs, y_: batch_ys, keep_prob: 0.5})
            train_avg_cost += c / total_batch
            c = sess.run(cost_drop, feed_dict={x: mnist.test.images, y_: mnist.
↪test.labels, keep_prob: 1.0})
            test_avg_cost += c / total_batch

            xs.append(epoch)
            ytrs.append(train_avg_cost)
            ytes.append(test_avg_cost)
            plt_dynamic(xs, ytrs, ytes, ax,np.arange(1, 1.8, step=0.05),
↪"input-ReLu(512)-Dropout-ReLu(128)-Dropout-Sigmoid(output)-AdamOptimizer")

            if epoch%display_step == 0:
                print("Epoch:", '%04d' % (epoch+1), "train cost={:.9f}".
↪format(train_avg_cost), "test cost={:.9f}".format(test_avg_cost))

        # plot final results
        plt_dynamic(xs, ytrs, ytes, ax,np.arange(1, 1.8, step=0.05),
↪"input-ReLu(512)-Dropout-ReLu(128)-Dropout-Sigmoid(output)-AdamOptimizer")

        # we are calculating the final accuracy on the test data
        correct_prediction = tf.equal(tf.argmax(ydrop,1), tf.argmax(y_,1))
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
        print("Accuracy:", accuracy.eval({x: mnist.test.images, y_: mnist.test.
↪labels, keep_prob: 1.0 })))

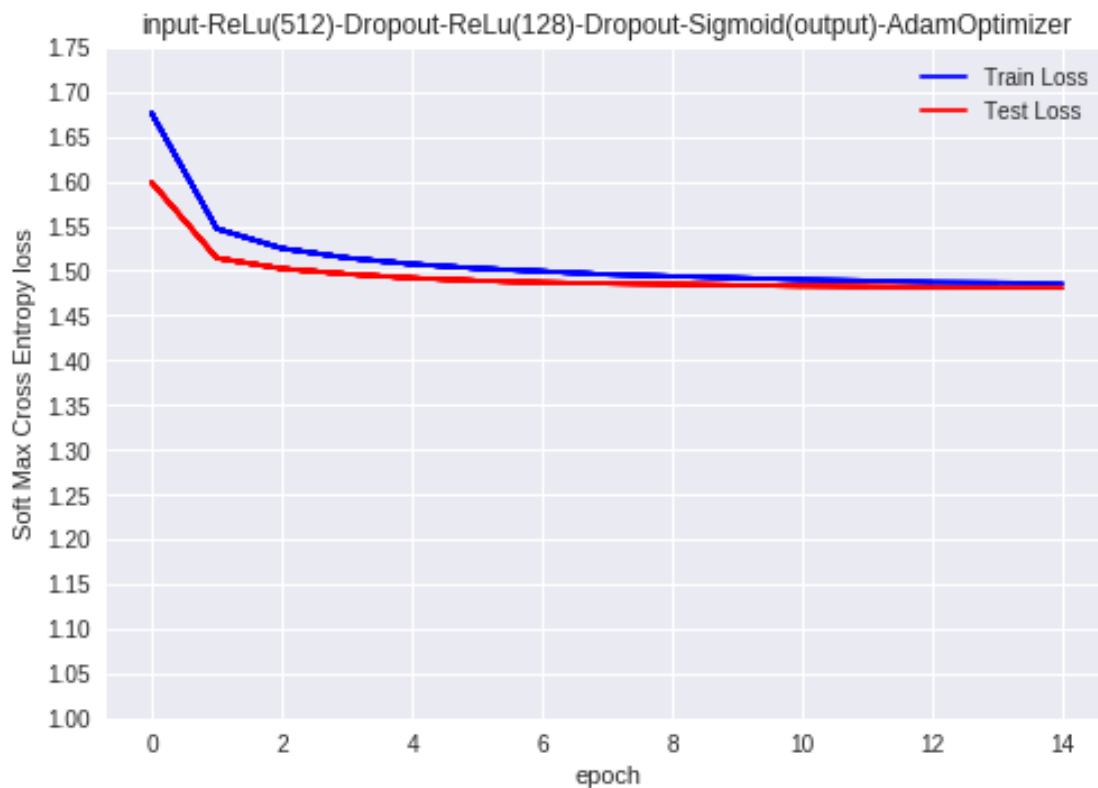
```

```

x: (?, 784) W[h1]: (784, 512) b[h1]: (512,)
layer_1: (?, 512) W[h2]: (512, 128) b[h2]: (128,)
layer_2: (?, 128) W[out]: (128, 10) b3: (10,)
out_layer: (?, 10)
Epoch: 0001 train cost=1.675481348 test cost=1.598583966

```

Epoch: 0002 train cost=1.546939785 test cost=1.514125155
Epoch: 0003 train cost=1.524996623 test cost=1.502378711
Epoch: 0004 train cost=1.514395935 test cost=1.496348261
Epoch: 0005 train cost=1.507464443 test cost=1.491901992
Epoch: 0006 train cost=1.502681051 test cost=1.488922666
Epoch: 0007 train cost=1.499395829 test cost=1.486873763
Epoch: 0008 train cost=1.495856912 test cost=1.485843970
Epoch: 0009 train cost=1.493913886 test cost=1.484971796
Epoch: 0010 train cost=1.491971510 test cost=1.483862281
Epoch: 0011 train cost=1.490074936 test cost=1.482910818
Epoch: 0012 train cost=1.488904240 test cost=1.482207586
Epoch: 0013 train cost=1.487357789 test cost=1.481696167
Epoch: 0014 train cost=1.486711373 test cost=1.480950590
Epoch: 0015 train cost=1.485781125 test cost=1.480801912
Accuracy: 0.9736



```
[0]: h1_w = w['h1'].flatten().reshape(-1,1)
      h2_w = w['h2'].flatten().reshape(-1,1)
      out_w = w['out'].flatten().reshape(-1,1)

      fig = plt.figure()
```

```

plt.subplot(1, 3, 1)
plt.title("Weight matrix")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Weight matrix ")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Weight matrix ")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
# plt.show()

```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)

```

[0]: Text(0.5,0,'Output Layer ')

