

# MNIST\_with\_Keras

September 16, 2020

## 0.1 Keras – MLPs on MNIST

```
[0]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow"
      ↪ use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

```
[0]: %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
[0]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 3s 0us/step

```
[0]: print("Number of training examples :", X_train.shape[0], "and each image is of_
      ↪ shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of_
      ↪ shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)

Number of training examples : 10000 and each image is of shape (28, 28)

```
[0]: # if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784
```

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
[0]: # after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of_
→shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of_
→shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)

Number of training examples : 10000 and each image is of shape (784)

```
[0]: # An example data point
```

```
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  3  18  18  18 126 136 175 26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0 30 36 94 154
170 253 253 253 253 253 225 172 253 242 195 64  0  0  0  0  0  0
 0  0  0  0  0 49 238 253 253 253 253 253 253 253 251 93 82
82 56 39  0  0  0  0  0  0  0  0  0  0  0  0 18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0 80 156 107 253 253 205 11  0 43 154
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0 14  1 154 253 90  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0 11 190 253 70  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0 81 240 253 253 119 25  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0 45 186 253 253 150 27  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0 16 93 252 253 187
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
```

0	0	0	0	0	0	0	249	253	249	64	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	46	130	183	253
253	207	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	39	148	229	253	253	253	250	182	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	24	114	221	253	253	253
253	201	78	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	23	66	213	253	253	253	253	198	81	2	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	18	171	219	253	253	253	253	195
80	9	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
55	172	226	253	253	253	253	244	133	11	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	136	253	253	253	212	135	132	16
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0]							

```
[0]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the
# data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

```
[0]: # example data point after normalizing
print(X_train[0])
```

[illegible]

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.01176471	0.07058824	0.07058824	0.07058824
0.49411765	0.53333333	0.68627451	0.10196078	0.65098039	1.
0.96862745	0.49803922	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.21568627	0.6745098	0.88627451	0.99215686	0.99215686	0.99215686
0.99215686	0.95686275	0.52156863	0.04313725	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.53333333	0.99215686
0.99215686	0.99215686	0.83137255	0.52941176	0.51764706	0.0627451
0.	0.	0.	0.	0.	0.

[illegible]

```
[0]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ", Y_train[0])
```

```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
Softmax classifier
```

```
[0]: # https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the
↳ constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:
```

```

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True,
    ↪ kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,
    ↪ activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) +
    ↪ bias) where
# activation is the element-wise activation function passed as the activation
    ↪ argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is
    ↪ True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the
    ↪ activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

```
[0]: # some model parameters
```

```
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

```
[0]: # start building a model
```

```
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

```
[0]: # Before training a model, you need to configure the learning process, which is
      ↪ done via the compile method
```

```
# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer ,
      ↪ https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize.,
      ↪ https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to
      ↪ metrics=['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in
      ↪ categorical format
# (e.g. if you have 10 classes, the target for each sample should be a
      ↪ 10-dimensional vector that is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted our labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy',
      ↪ metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
```



```

# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1,
    ↳callbacks=None, validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None,
    ↳initial_epoch=0, steps_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a
    ↳dataset).

# it returns A History object. Its History.history attribute is a record of
    ↳training loss values and
# metrics values at successive epochs, as well as validation loss values and
    ↳validation metrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
    ↳verbose=1, validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 49us/step - loss: 1.2935 -  
acc: 0.6829 - val\_loss: 0.8171 - val\_acc: 0.8312

Epoch 2/20

60000/60000 [=====] - 2s 35us/step - loss: 0.7213 -  
acc: 0.8385 - val\_loss: 0.6099 - val\_acc: 0.8623

Epoch 3/20

60000/60000 [=====] - 2s 35us/step - loss: 0.5904 -  
acc: 0.8584 - val\_loss: 0.5275 - val\_acc: 0.8741

Epoch 4/20

60000/60000 [=====] - 2s 35us/step - loss: 0.5278 -  
acc: 0.8682 - val\_loss: 0.4815 - val\_acc: 0.8814

Epoch 5/20

60000/60000 [=====] - 2s 35us/step - loss: 0.4897 -  
acc: 0.8747 - val\_loss: 0.4515 - val\_acc: 0.8870

Epoch 6/20

34432/60000 [=====>...] - ETA: 0s - loss: 0.4697 - acc:  
0.877360000/60000 [=====] - 2s 35us/step - loss: 0.4635  
- acc: 0.8799 - val\_loss: 0.4303 - val\_acc: 0.8898

Epoch 7/20

60000/60000 [=====] - 2s 35us/step - loss: 0.4442 -  
acc: 0.8837 - val\_loss: 0.4138 - val\_acc: 0.8917

Epoch 8/20

60000/60000 [=====] - 2s 35us/step - loss: 0.4290 -  
acc: 0.8866 - val\_loss: 0.4010 - val\_acc: 0.8952

```

Epoch 9/20
60000/60000 [=====] - 2s 35us/step - loss: 0.4169 -
acc: 0.8894 - val_loss: 0.3909 - val_acc: 0.8971
Epoch 10/20
60000/60000 [=====] - 2s 35us/step - loss: 0.4067 -
acc: 0.8911 - val_loss: 0.3818 - val_acc: 0.8994
Epoch 11/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3982 -
acc: 0.8926 - val_loss: 0.3743 - val_acc: 0.9006
Epoch 12/20
1792/60000 [...] - ETA: 1s - loss: 0.4077 - acc:
0.889560000/60000 [=====] - 2s 35us/step - loss: 0.3908
- acc: 0.8948 - val_loss: 0.3681 - val_acc: 0.9020
Epoch 13/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3844 -
acc: 0.8960 - val_loss: 0.3624 - val_acc: 0.9039
Epoch 14/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3786 -
acc: 0.8972 - val_loss: 0.3575 - val_acc: 0.9046
Epoch 15/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3735 -
acc: 0.8981 - val_loss: 0.3528 - val_acc: 0.9058
Epoch 16/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3689 -
acc: 0.8993 - val_loss: 0.3490 - val_acc: 0.9062
Epoch 17/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3648 -
acc: 0.9004 - val_loss: 0.3455 - val_acc: 0.9066
Epoch 18/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3610 -
acc: 0.9016 - val_loss: 0.3419 - val_acc: 0.9077
Epoch 19/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3575 -
acc: 0.9024 - val_loss: 0.3389 - val_acc: 0.9088
Epoch 20/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3544 -
acc: 0.9032 - val_loss: 0.3362 - val_acc: 0.9093

```

```

[0]: score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

```

```

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
    ↳ epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter
    ↳ validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
    ↳ number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.3362289469957352

Test accuracy: 0.9093

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

MLP + Sigmoid activation + SGDOptimizer

```

[0]: # Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

```

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 512)	401920
dense_3 (Dense)	(None, 128)	65664
dense_4 (Dense)	(None, 10)	1290

Total params: 468,874  
Trainable params: 468,874  
Non-trainable params: 0

```
[0]: model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy',  
    ↪metrics=['accuracy'])  
  
history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size,  
    ↪epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 42us/step - loss: 2.2708 -  
acc: 0.2169 - val\_loss: 2.2197 - val\_acc: 0.2768

Epoch 2/20

60000/60000 [=====] - 2s 42us/step - loss: 2.1739 -  
acc: 0.4237 - val\_loss: 2.1143 - val\_acc: 0.4877

Epoch 3/20

60000/60000 [=====] - 2s 42us/step - loss: 2.0506 -  
acc: 0.5485 - val\_loss: 1.9659 - val\_acc: 0.5524

Epoch 4/20

60000/60000 [=====] - 3s 42us/step - loss: 1.8796 -  
acc: 0.6135 - val\_loss: 1.7673 - val\_acc: 0.6481

Epoch 5/20

60000/60000 [=====] - 3s 42us/step - loss: 1.6674 -  
acc: 0.6659 - val\_loss: 1.5414 - val\_acc: 0.7205

Epoch 6/20

5376/60000 [=>...] - ETA: 2s - loss: 1.5430 - acc:  
0.698160000/60000 [=====] - 2s 41us/step - loss: 1.4449  
- acc: 0.7125 - val\_loss: 1.3233 - val\_acc: 0.7513

Epoch 7/20

60000/60000 [=====] - 2s 41us/step - loss: 1.2442 -  
acc: 0.7466 - val\_loss: 1.1406 - val\_acc: 0.7599

Epoch 8/20

60000/60000 [=====] - 2s 41us/step - loss: 1.0815 -  
acc: 0.7722 - val\_loss: 0.9974 - val\_acc: 0.7968

Epoch 9/20

60000/60000 [=====] - 2s 41us/step - loss: 0.9544 -  
acc: 0.7940 - val\_loss: 0.8867 - val\_acc: 0.8060

Epoch 10/20

60000/60000 [=====] - 2s 41us/step - loss: 0.8556 -  
acc: 0.8082 - val\_loss: 0.7984 - val\_acc: 0.8227

Epoch 11/20

31232/60000 [=====>...] - ETA: 1s - loss: 0.7920 - acc:  
0.820760000/60000 [=====] - 2s 42us/step - loss: 0.7776  
- acc: 0.8225 - val\_loss: 0.7292 - val\_acc: 0.8335

Epoch 12/20

```

60000/60000 [=====] - 2s 41us/step - loss: 0.7154 -
acc: 0.8333 - val_loss: 0.6741 - val_acc: 0.8422
Epoch 13/20
60000/60000 [=====] - 2s 41us/step - loss: 0.6650 -
acc: 0.8412 - val_loss: 0.6282 - val_acc: 0.8500
Epoch 14/20
60000/60000 [=====] - 2s 41us/step - loss: 0.6237 -
acc: 0.8485 - val_loss: 0.5906 - val_acc: 0.8585
Epoch 15/20
60000/60000 [=====] - 2s 41us/step - loss: 0.5893 -
acc: 0.8546 - val_loss: 0.5591 - val_acc: 0.8616
Epoch 16/20
34432/60000 [=====>...] - ETA: 0s - loss: 0.5691 - acc:
0.857860000/60000 [=====] - 2s 41us/step - loss: 0.5606
- acc: 0.8599 - val_loss: 0.5329 - val_acc: 0.8672
Epoch 17/20
60000/60000 [=====] - 2s 41us/step - loss: 0.5361 -
acc: 0.8641 - val_loss: 0.5095 - val_acc: 0.8703
Epoch 18/20
60000/60000 [=====] - 2s 42us/step - loss: 0.5151 -
acc: 0.8676 - val_loss: 0.4904 - val_acc: 0.8736
Epoch 19/20
60000/60000 [=====] - 2s 41us/step - loss: 0.4969 -
acc: 0.8710 - val_loss: 0.4732 - val_acc: 0.8782
Epoch 20/20
60000/60000 [=====] - 2s 42us/step - loss: 0.4810 -
acc: 0.8739 - val_loss: 0.4583 - val_acc: 0.8816

```

```

[0]: score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
→ epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter
→ validation_data
# val_loss : validation loss
# val_acc : validation accuracy

```

```

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
    ↳ number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.4582893396139145

Test accuracy: 0.8816

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```

[0]: w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)

```

MLP + Sigmoid activation + ADAM

```

[0]: model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size,
    ↪epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

```

-----
Layer (type)                Output Shape                Param #
=====
dense_5 (Dense)              (None, 512)                 401920
-----
dense_6 (Dense)              (None, 128)                 65664
-----
dense_7 (Dense)              (None, 10)                  1290
=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0
-----
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 3s 55us/step - loss: 0.5308 -
acc: 0.8636 - val_loss: 0.2550 - val_acc: 0.9259
Epoch 2/20
53120/60000 [=====>...] - ETA: 0s - loss: 0.2244 - acc:
0.934060000/60000 [=====] - 3s 51us/step - loss: 0.2205
- acc: 0.9351 - val_loss: 0.1946 - val_acc: 0.9417
Epoch 3/20
60000/60000 [=====] - 3s 51us/step - loss: 0.1650 -
acc: 0.9512 - val_loss: 0.1421 - val_acc: 0.9570
Epoch 4/20
60000/60000 [=====] - 3s 51us/step - loss: 0.1279 -

```

```

acc: 0.9614 - val_loss: 0.1238 - val_acc: 0.9645
Epoch 5/20
60000/60000 [=====] - 3s 51us/step - loss: 0.1010 -
acc: 0.9704 - val_loss: 0.1029 - val_acc: 0.9693
Epoch 6/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0801 -
acc: 0.9763 - val_loss: 0.0877 - val_acc: 0.9725
Epoch 7/20
 4480/60000 [=>...] - ETA: 2s - loss: 0.0632 - acc:
0.979560000/60000 [=====] - 3s 51us/step - loss: 0.0645
- acc: 0.9809 - val_loss: 0.0831 - val_acc: 0.9751
Epoch 8/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0519 -
acc: 0.9842 - val_loss: 0.0724 - val_acc: 0.9780
Epoch 9/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0433 -
acc: 0.9872 - val_loss: 0.0714 - val_acc: 0.9786
Epoch 10/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0347 -
acc: 0.9898 - val_loss: 0.0695 - val_acc: 0.9776
Epoch 11/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0268 -
acc: 0.9930 - val_loss: 0.0659 - val_acc: 0.9796
Epoch 12/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0219 -
acc: 0.9944 - val_loss: 0.0642 - val_acc: 0.9809
Epoch 13/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0180 -
acc: 0.9953 - val_loss: 0.0677 - val_acc: 0.9794
Epoch 14/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0133 -
acc: 0.9970 - val_loss: 0.0647 - val_acc: 0.9803
Epoch 15/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0114 -
acc: 0.9975 - val_loss: 0.0628 - val_acc: 0.9812
Epoch 16/20
57344/60000 [=====>...] - ETA: 0s - loss: 0.0085 - acc:
0.998260000/60000 [=====] - 3s 50us/step - loss: 0.0085
- acc: 0.9982 - val_loss: 0.0666 - val_acc: 0.9806
Epoch 17/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0070 -
acc: 0.9986 - val_loss: 0.0643 - val_acc: 0.9822
Epoch 18/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0061 -
acc: 0.9986 - val_loss: 0.0656 - val_acc: 0.9818
Epoch 19/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0055 -
acc: 0.9988 - val_loss: 0.0811 - val_acc: 0.9774

```



Epoch 20/20  
60000/60000 [=====] - 3s 50us/step - loss: 0.0038 -  
acc: 0.9992 - val\_loss: 0.0723 - val\_acc: 0.9818

```
[0]: score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
#     ↪ epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter
#     ↪ validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
#     ↪ number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06385514608082886

Test accuracy: 0.9824

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```
[0]: w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816:
FutureWarning: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)

```

MLP + ReLU +SGD

```

[0]: # Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0, )$  we satisfy this
→ condition with  $=\sqrt{2/(n_i)}$ .
# h1 =>  $=\sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, ) = N(0, 0.062)$ 
# h2 =>  $=\sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, ) = N(0, 0.125)$ 
# out =>  $=\sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0, ) = N(0, 0.120)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,),
→kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu',
→kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

```

```
model_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dense_10 (Dense)	(None, 10)	1290

Total params: 468,874  
Trainable params: 468,874  
Non-trainable params: 0

```
[0]: model_relu.compile(optimizer='sgd', loss='categorical_crossentropy',  
    ↪ metrics=['accuracy'])  
  
history = model_relu.fit(X_train, Y_train, batch_size=batch_size,  
    ↪ epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 67us/step - loss: 0.7579 -  
acc: 0.7812 - val\_loss: 0.3951 - val\_acc: 0.8921

Epoch 2/20

60000/60000 [=====] - 4s 64us/step - loss: 0.3535 -  
acc: 0.8998 - val\_loss: 0.3040 - val\_acc: 0.9153

Epoch 3/20

60000/60000 [=====] - 4s 64us/step - loss: 0.2900 -  
acc: 0.9172 - val\_loss: 0.2648 - val\_acc: 0.9253

Epoch 4/20

60000/60000 [=====] - 4s 60us/step - loss: 0.2558 -  
acc: 0.9269 - val\_loss: 0.2393 - val\_acc: 0.9316

Epoch 5/20

60000/60000 [=====] - 4s 58us/step - loss: 0.2324 -  
acc: 0.9340 - val\_loss: 0.2210 - val\_acc: 0.9371

Epoch 6/20

60000/60000 [=====] - 4s 64us/step - loss: 0.2144 -  
acc: 0.9391 - val\_loss: 0.2072 - val\_acc: 0.9400

Epoch 7/20

60000/60000 [=====] - 4s 66us/step - loss: 0.1995 -  
acc: 0.9443 - val\_loss: 0.1957 - val\_acc: 0.9444

Epoch 8/20

60000/60000 [=====] - 4s 61us/step - loss: 0.1872 -

```

acc: 0.9476 - val_loss: 0.1848 - val_acc: 0.9456
Epoch 9/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1763 -
acc: 0.9507 - val_loss: 0.1771 - val_acc: 0.9488
Epoch 10/20
60000/60000 [=====] - 4s 60us/step - loss: 0.1668 -
acc: 0.9539 - val_loss: 0.1682 - val_acc: 0.9506
Epoch 11/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1585 -
acc: 0.9560 - val_loss: 0.1623 - val_acc: 0.9518
Epoch 12/20
60000/60000 [=====] - 7s 113us/step - loss: 0.1511 -
acc: 0.9577 - val_loss: 0.1560 - val_acc: 0.9543
Epoch 13/20
60000/60000 [=====] - 7s 115us/step - loss: 0.1443 -
acc: 0.9596 - val_loss: 0.1517 - val_acc: 0.9557
Epoch 14/20
60000/60000 [=====] - 7s 111us/step - loss: 0.1379 -
acc: 0.9615 - val_loss: 0.1474 - val_acc: 0.9572
Epoch 15/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1323 -
acc: 0.9628 - val_loss: 0.1429 - val_acc: 0.9580
Epoch 16/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1270 -
acc: 0.9645 - val_loss: 0.1371 - val_acc: 0.9598
Epoch 17/20
60000/60000 [=====] - 7s 110us/step - loss: 0.1221 -
acc: 0.9661 - val_loss: 0.1351 - val_acc: 0.9602
Epoch 18/20
60000/60000 [=====] - 4s 62us/step - loss: 0.1177 -
acc: 0.9671 - val_loss: 0.1309 - val_acc: 0.9618
Epoch 19/20
60000/60000 [=====] - 4s 60us/step - loss: 0.1136 -
acc: 0.9685 - val_loss: 0.1263 - val_acc: 0.9631
Epoch 20/20
60000/60000 [=====] - 5s 79us/step - loss: 0.1094 -
acc: 0.9694 - val_loss: 0.1241 - val_acc: 0.9631

```

```

[0]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

```

```

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
    ↳ epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter
    ↳ validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
    ↳ number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.12405014228336513

Test accuracy: 0.9631

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```

[0]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

MLP + ReLU + ADAM

```
[0]: model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,),
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu',
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy',
    ↳metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size,
    ↳epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
-----
Layer (type)                Output Shape                Param #
=====
dense_11 (Dense)            (None, 512)                 401920
-----
dense_12 (Dense)            (None, 128)                 65664
-----
dense_13 (Dense)            (None, 10)                  1290
=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0
-----
None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 121us/step - loss: 0.2341 -
acc: 0.9295 - val_loss: 0.1165 - val_acc: 0.9652
Epoch 2/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0878 -
```

acc: 0.9729 - val\_loss: 0.0883 - val\_acc: 0.9720  
 Epoch 3/20  
 60000/60000 [=====] - 5s 75us/step - loss: 0.0544 -  
 acc: 0.9825 - val\_loss: 0.0860 - val\_acc: 0.9729  
 Epoch 4/20  
 60000/60000 [=====] - 4s 70us/step - loss: 0.0354 -  
 acc: 0.9885 - val\_loss: 0.0699 - val\_acc: 0.9797  
 Epoch 5/20  
 60000/60000 [=====] - 4s 73us/step - loss: 0.0266 -  
 acc: 0.9914 - val\_loss: 0.0720 - val\_acc: 0.9788  
 Epoch 6/20  
 60000/60000 [=====] - 4s 70us/step - loss: 0.0200 -  
 acc: 0.9941 - val\_loss: 0.0696 - val\_acc: 0.9803  
 Epoch 7/20  
 60000/60000 [=====] - 4s 73us/step - loss: 0.0155 -  
 acc: 0.9951 - val\_loss: 0.0640 - val\_acc: 0.9829  
 Epoch 8/20  
 60000/60000 [=====] - 4s 71us/step - loss: 0.0140 -  
 acc: 0.9952 - val\_loss: 0.0848 - val\_acc: 0.9792  
 Epoch 9/20  
 60000/60000 [=====] - 4s 71us/step - loss: 0.0143 -  
 acc: 0.9952 - val\_loss: 0.0837 - val\_acc: 0.9796  
 Epoch 10/20  
 60000/60000 [=====] - 7s 115us/step - loss: 0.0128 -  
 acc: 0.9958 - val\_loss: 0.0946 - val\_acc: 0.9782  
 Epoch 11/20  
 60000/60000 [=====] - 7s 125us/step - loss: 0.0081 -  
 acc: 0.9974 - val\_loss: 0.0682 - val\_acc: 0.9826  
 Epoch 12/20  
 60000/60000 [=====] - 8s 129us/step - loss: 0.0121 -  
 acc: 0.9959 - val\_loss: 0.0793 - val\_acc: 0.9816  
 Epoch 13/20  
 60000/60000 [=====] - 8s 133us/step - loss: 0.0107 -  
 acc: 0.9963 - val\_loss: 0.0746 - val\_acc: 0.9820  
 Epoch 14/20  
 60000/60000 [=====] - 8s 129us/step - loss: 0.0113 -  
 acc: 0.9960 - val\_loss: 0.0813 - val\_acc: 0.9816  
 Epoch 15/20  
 60000/60000 [=====] - 5s 77us/step - loss: 0.0058 -  
 acc: 0.9982 - val\_loss: 0.0770 - val\_acc: 0.9842  
 Epoch 16/20  
 60000/60000 [=====] - 4s 65us/step - loss: 0.0040 -  
 acc: 0.9987 - val\_loss: 0.0930 - val\_acc: 0.9808  
 Epoch 17/20  
 60000/60000 [=====] - 4s 68us/step - loss: 0.0119 -  
 acc: 0.9959 - val\_loss: 0.0813 - val\_acc: 0.9819  
 Epoch 18/20  
 60000/60000 [=====] - 4s 73us/step - loss: 0.0105 -

```

acc: 0.9966 - val_loss: 0.1000 - val_acc: 0.9803
Epoch 19/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0064 -
acc: 0.9981 - val_loss: 0.0852 - val_acc: 0.9831
Epoch 20/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0056 -
acc: 0.9982 - val_loss: 0.1029 - val_acc: 0.9805

```

```

[0]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
    ↳ epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter
    ↳ validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
    ↳ number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10294274219236926

Test accuracy: 0.9805

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>



```
[0]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

```
[0]: # Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0, )$  we satisfy this
# condition with  $=\sqrt{2/(n_i+n_i+1)}$ .
# h1 =>  $=\sqrt{2/(n_i+n_i+1)} = 0.039 \Rightarrow N(0, ) = N(0, 0.039)$ 
# h2 =>  $=\sqrt{2/(n_i+n_i+1)} = 0.055 \Rightarrow N(0, ) = N(0, 0.055)$ 
# h1 =>  $=\sqrt{2/(n_i+n_i+1)} = 0.120 \Rightarrow N(0, ) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,),
# kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
```

```

model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid',
    ↳kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

```

-----
Layer (type)                Output Shape          Param #
=====
dense_14 (Dense)             (None, 512)           401920
-----
batch_normalization_1 (Batch Normalization) (None, 512)           2048
-----
dense_15 (Dense)             (None, 128)           65664
-----
batch_normalization_2 (Batch Normalization) (None, 128)           512
-----
dense_16 (Dense)             (None, 10)            1290
=====
Total params: 471,434
Trainable params: 470,154
Non-trainable params: 1,280
-----

```

```

[0]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy',
    ↳metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size,
    ↳epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 138us/step - loss: 0.3036 -  
acc: 0.9104 - val\_loss: 0.2116 - val\_acc: 0.9376

Epoch 2/20

60000/60000 [=====] - 10s 170us/step - loss: 0.1747 -  
acc: 0.9483 - val\_loss: 0.1670 - val\_acc: 0.9505

Epoch 3/20

60000/60000 [=====] - 13s 220us/step - loss: 0.1367 -  
acc: 0.9599 - val\_loss: 0.1451 - val\_acc: 0.9567

Epoch 4/20

60000/60000 [=====] - 9s 156us/step - loss: 0.1134 -

acc: 0.9666 - val\_loss: 0.1335 - val\_acc: 0.9603  
Epoch 5/20  
60000/60000 [=====] - 13s 211us/step - loss: 0.0949 -  
acc: 0.9703 - val\_loss: 0.1325 - val\_acc: 0.9589  
Epoch 6/20  
60000/60000 [=====] - 7s 119us/step - loss: 0.0802 -  
acc: 0.9758 - val\_loss: 0.1139 - val\_acc: 0.9652  
Epoch 7/20  
60000/60000 [=====] - 8s 127us/step - loss: 0.0682 -  
acc: 0.9787 - val\_loss: 0.1136 - val\_acc: 0.9666  
Epoch 8/20  
60000/60000 [=====] - 7s 124us/step - loss: 0.0608 -  
acc: 0.9815 - val\_loss: 0.1114 - val\_acc: 0.9666  
Epoch 9/20  
60000/60000 [=====] - 8s 129us/step - loss: 0.0532 -  
acc: 0.9837 - val\_loss: 0.1167 - val\_acc: 0.9666  
Epoch 10/20  
60000/60000 [=====] - 7s 123us/step - loss: 0.0455 -  
acc: 0.9856 - val\_loss: 0.0962 - val\_acc: 0.9718  
Epoch 11/20  
60000/60000 [=====] - 7s 112us/step - loss: 0.0376 -  
acc: 0.9880 - val\_loss: 0.1102 - val\_acc: 0.9673  
Epoch 12/20  
60000/60000 [=====] - 7s 124us/step - loss: 0.0350 -  
acc: 0.9889 - val\_loss: 0.1033 - val\_acc: 0.9710  
Epoch 13/20  
60000/60000 [=====] - 7s 124us/step - loss: 0.0308 -  
acc: 0.9903 - val\_loss: 0.1020 - val\_acc: 0.9712  
Epoch 14/20  
60000/60000 [=====] - 7s 123us/step - loss: 0.0271 -  
acc: 0.9913 - val\_loss: 0.1038 - val\_acc: 0.9727  
Epoch 15/20  
60000/60000 [=====] - 7s 122us/step - loss: 0.0231 -  
acc: 0.9926 - val\_loss: 0.1019 - val\_acc: 0.9717  
Epoch 16/20  
60000/60000 [=====] - 8s 127us/step - loss: 0.0220 -  
acc: 0.9928 - val\_loss: 0.1110 - val\_acc: 0.9703  
Epoch 17/20  
60000/60000 [=====] - 7s 114us/step - loss: 0.0229 -  
acc: 0.9928 - val\_loss: 0.1067 - val\_acc: 0.9739  
Epoch 18/20  
60000/60000 [=====] - 8s 128us/step - loss: 0.0203 -  
acc: 0.9935 - val\_loss: 0.0982 - val\_acc: 0.9738  
Epoch 19/20  
60000/60000 [=====] - 7s 125us/step - loss: 0.0171 -  
acc: 0.9944 - val\_loss: 0.1056 - val\_acc: 0.9706  
Epoch 20/20  
60000/60000 [=====] - 11s 182us/step - loss: 0.0146 -

acc: 0.9952 - val\_loss: 0.1046 - val\_acc: 0.9732

```
[0]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
#     → epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter
#     → validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
#     → number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10456635547156475

Test accuracy: 0.9732

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```
[0]: w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
```

```
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

## 5. MLP + Dropout + AdamOptimizer

```
[0]: # https://stackoverflow.com/questions/34716454/
      ↪ where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,),
    ↪kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid',
    ↪kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

-----  
Layer (type)

Output Shape

Param #

```

=====
dense_17 (Dense)                (None, 512)                401920
-----
batch_normalization_3 (Batch Normalization) (None, 512)                2048
-----
dropout_1 (Dropout)             (None, 512)                 0
-----
dense_18 (Dense)                (None, 128)               65664
-----
batch_normalization_4 (Batch Normalization) (None, 128)                512
-----
dropout_2 (Dropout)             (None, 128)                 0
-----
dense_19 (Dense)                (None, 10)                 1290
=====
Total params: 471,434
Trainable params: 470,154
Non-trainable params: 1,280
-----

```

```

[0]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪ metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
    ↪ epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 14s 227us/step - loss: 0.6612 -
acc: 0.7951 - val_loss: 0.2860 - val_acc: 0.9166
Epoch 2/20
60000/60000 [=====] - 8s 136us/step - loss: 0.4250 -
acc: 0.8710 - val_loss: 0.2545 - val_acc: 0.9252
Epoch 3/20
60000/60000 [=====] - 12s 198us/step - loss: 0.3841 -
acc: 0.8846 - val_loss: 0.2391 - val_acc: 0.9298
Epoch 4/20
60000/60000 [=====] - 8s 138us/step - loss: 0.3551 -
acc: 0.8927 - val_loss: 0.2279 - val_acc: 0.9325
Epoch 5/20
60000/60000 [=====] - 7s 123us/step - loss: 0.3355 -
acc: 0.8986 - val_loss: 0.2127 - val_acc: 0.9356
Epoch 6/20
60000/60000 [=====] - 8s 136us/step - loss: 0.3234 -
acc: 0.9031 - val_loss: 0.2029 - val_acc: 0.9387: 1s - loss:
Epoch 7/20
60000/60000 [=====] - 8s 131us/step - loss: 0.3068 -
acc: 0.9077 - val_loss: 0.1927 - val_acc: 0.9421

```

```

Epoch 8/20
60000/60000 [=====] - 11s 185us/step - loss: 0.2933 -
acc: 0.9113 - val_loss: 0.1836 - val_acc: 0.9453
Epoch 9/20
60000/60000 [=====] - 13s 222us/step - loss: 0.2850 -
acc: 0.9131 - val_loss: 0.1797 - val_acc: 0.9451
Epoch 10/20
60000/60000 [=====] - 14s 236us/step - loss: 0.2715 -
acc: 0.9187 - val_loss: 0.1738 - val_acc: 0.9465
Epoch 11/20
60000/60000 [=====] - 8s 141us/step - loss: 0.2611 -
acc: 0.9214 - val_loss: 0.1671 - val_acc: 0.9506
Epoch 12/20
60000/60000 [=====] - 8s 134us/step - loss: 0.2464 -
acc: 0.9252 - val_loss: 0.1554 - val_acc: 0.9525
Epoch 13/20
60000/60000 [=====] - 8s 137us/step - loss: 0.2382 -
acc: 0.9278 - val_loss: 0.1479 - val_acc: 0.9554
Epoch 14/20
60000/60000 [=====] - 8s 136us/step - loss: 0.2275 -
acc: 0.9313 - val_loss: 0.1375 - val_acc: 0.9580
Epoch 15/20
60000/60000 [=====] - 8s 137us/step - loss: 0.2183 -
acc: 0.9337 - val_loss: 0.1326 - val_acc: 0.9599
Epoch 16/20
60000/60000 [=====] - 8s 138us/step - loss: 0.2068 -
acc: 0.9384 - val_loss: 0.1297 - val_acc: 0.9613 loss: 0.2066 - ac
Epoch 17/20
60000/60000 [=====] - 8s 139us/step - loss: 0.2011 -
acc: 0.9395 - val_loss: 0.1181 - val_acc: 0.9646
Epoch 18/20
60000/60000 [=====] - 8s 137us/step - loss: 0.1886 -
acc: 0.9435 - val_loss: 0.1145 - val_acc: 0.9658
Epoch 19/20
60000/60000 [=====] - 8s 138us/step - loss: 0.1821 -
acc: 0.9451 - val_loss: 0.1104 - val_acc: 0.9662
Epoch 20/20
60000/60000 [=====] - 8s 139us/step - loss: 0.1739 -
acc: 0.9473 - val_loss: 0.1093 - val_acc: 0.9679

```

```

[0]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

```

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

```

```

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
    ↳ epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter
    ↳ validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to
    ↳ number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.1093290721397847

Test accuracy: 0.9679

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```

[0]: w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')

```



```
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

Hyper-parameter tuning of Keras models using Sklearn

```
[0]: from keras.optimizers import Adam,RMSprop,SGD
def best_hyperparameters(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,),
↪kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(Dense(128, activation=activ,
↪kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'],
↪optimizer='adam')

    return model
```

```
[0]: # https://machinelearningmastery.com/
↪grid-search-hyperparameters-deep-learning-models-python-keras/

activ = ['sigmoid','relu']

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV

model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch,
↪batch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter
```

```
grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)
```

```
[0]: print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
      means = grid_result.cv_results_['mean_test_score']
      stds = grid_result.cv_results_['std_test_score']
      params = grid_result.cv_results_['params']
      for mean, stdev, param in zip(means, stds, params):
          print("%f (%f) with: %r" % (mean, stdev, param))
```

```
Best: 0.975633 using {'activ': 'relu'}
0.974650 (0.001138) with: {'activ': 'sigmoid'}
0.975633 (0.002812) with: {'activ': 'relu'}
```